



Tracing Borrowed Authority in Relational AI Systems

AIES 2026 | STUDENT PROGRAM

Hanjing Shi Lehigh University hasa23@lehigh.edu

TL;DR A familiar identity does not guarantee sound judgment or accountability.

Persona AI presents generated responses through a recognizable identity. That identity can lend an answer weight without supplying the judgment, permission or support behind it. Across experts, family/friends and pets, we examine who supplies, checks, authorizes and can challenge the output.

01 Introduction | Why identity matters

Öhman’s The Afterlife of Data asks who controls our digital remains [4]; earlier work already framed afterlives as a governance issue [5]. Generative AI adds new speech [6]. Tool-using agents may connect that speech to decisions and action. A familiar identity can persist while its model, sources and operator change.

Who gives the response weight—and who does the work behind it?

02 Related work | From source to responsibility

Digital-afterlife, advice-taking and mimetic-AI research address different parts of the problem. Our synthesis of 50 scholarly sources and 8 public artifacts connects them [2].

Whose identity? → Why listen? → What reliance? → Who answers?

This is an analytic framework; the full sequence remains untested.

03 One question across three identity contexts

A Expert persona

Reputation and domain knowledge

A familiar expert name can make advice worth hearing.

Who provides current evidence, checks accuracy and judges fit?

Expert identity does not itself supply expert work.

B Family / friend persona

Shared history and personal knowledge

A familiar voice may feel like someone who knows me.

Which words are recorded, inferred or invented—and who may authorize them?

An archive cannot settle a new judgment.

C Companion-animal persona

Care memories and emotional bonds

A particular pet can make an encounter meaningful.

What may this representation say, and who sets its limits?

Comfort does not imply advisory competence.

Across all three: Who authorizes the representation, supplies and checks its outputs, and can correct or stop it?

These contexts can overlap. They are comparative lenses, not audit categories or tested experimental conditions.

04 Methods | Public-service audit [1]

Web, app-store, directory and literature searches yielded 93 records. One coder read public product pages, FAQs and policies before sign-up or payment. Fields covered authorization, source grounding, redress, deletion/export and support boundaries. No conversations or grief outcomes were tested.

29 Human afterlife services

Voices, stories, likenesses or conversations: permission to represent a person and generate new attributed speech.

24 Static human memorial services

Pages and bounded media, sometimes AI-generated: stewardship and exit without an ongoing persona.

17 Companion-animal services (memorials to personas)

Emotional bonds support inclusion [8]; shared care and inability to consent require separate analysis [1].

70 memorial-facing records + 23 adjacent tools (17 persona/companion; 6 mimetic). Analysis retains Human 53 / Animal 17 / Adjacent 23. These sampling families differ from the identity contexts above.

06 Situated evidence | What each identity leaves unresolved

A Expert | Zhang Xuefeng

Zhang Xuefeng is a Chinese education advisor whose unofficial posthumous personas offer admissions advice. We thematically analyzed 115 focal Xiaohongshu posts and 917 text-bearing comments/replies across 53 threads [3].

An installation ad promised expert consultation but sold setup. A separate family account added nearly 600 scanned pages of admissions records, checked errors and assessed student fit; the student made the final choice. Posts also disputed authorization to speak under Zhang’s name or keep assigning the deceased new work.

Finding: expert identity leaves updating and judgment to living actors. Reported use does not verify admissions outcomes.

B Family / friend | Public cases

Public cases in the afterlife study [1] compare how a person or loved one supplies source material. These are documented public accounts, not interviews with families; friendship-specific reliance remains untested.

Dadbot linked a father’s recorded stories to authored responses, with the father’s participation while alive. Project December generated new replies from a deceased partner’s messages and a prompt; the reviewed record did not establish her prior authorization.

Access to intimate records can enable reconstruction without settling permission for new speech or whose interpretation prevails.

Implication: distinguish recorded from generated speech and let affected relatives question sources, authority and continuation.

C Companion animal | Afterglow

Afterglow is a place-based pet-memorial prototype with private memories, generated pet messages and shared remembrance. A survey (N=57) and six roundtables with 20 unique adults informed iterative design [7].

In prototype discussions, participants imagined returning to share everyday joys and difficulties with their pet. One participant rejected life advice because it felt like a chatbot pretending to be their cat; another rejected a therapist-like role. The pet’s likeness can invite connection while also making some generated responses unacceptable.

Interpretation: meaningful companionship need not grant advisory authority. These were anticipated uses, not long-term deployment outcomes.

07 Contributions | Trace authority, work and responsibility

Conceptual

Relational provenance separates whose identity is invoked from who supplies the output. The framework connects potential authority, reliance and later accountability [2].

Limits: nonrepresentative public records and formative design encounters; comments are contributions, not independent people. No matched three-group comparison, causal reliance effect, long-term dependency outcome or autonomous harm is established.

Governance

Make authorization, sources, updating responsibilities, behavioral limits, objection and exit visible. Keep professional verification and relational care explicit.

Next study — planned

An adult vignette study will hold advice constant and vary persona source and responsibility labels. Measure advice acceptance, legitimacy, credit/blame and contestability separately.

[1] Shi & DiFranzo (2026). Beloved Afterlives. [arXiv:2609.29544](#) (preprint).

[2] Shi (2026). When AI Advice Speaks Through Someone Else. Manuscript.

[3] His Name, Their Judgment: Expert Authority in Posthumous Persona AI. Manuscript (2026).

[4] Öhman (2024). The Afterlife of Data. University of Chicago Press.

[5] Öhman & Floridi (2018). An ethical framework for the digital afterlife industry. NHB.

[6] Morris & Brubaker (2025). Generative Ghosts. CHI.

[7] Shi & DiFranzo (2026). Afterglow: A Place-Based Memorial Ecology. Manuscript.

[8] Bowen et al. (2021). Companion dogs as social support. doi:10.3389/fpsy.2021.622060.

